

Introduction of Mesos persistent storage

Weitao Zhou @ DataMan

Content

- How to run stateful service against current Mesos-0.22
- Disk isolation and monitoring
- Persistent Volumes
- Dynamic Reservations
- What we can contribute for Mesos persistent storage

How to run stateful service against current Mesos-0.22

- Making it outside of Mesos cluster
- Storing data into local filesystem of the specified Mesos slave
- Setting DFS to get rid of static reserve
- Taking advantage of stateful service build-in data-replication feature

Making it outside of Mesos Cluster

- Scenarios:
 - independent cluster: RabbitMQ cluster
 - constant resource usage, etc.
- Pros:
 - mature solution
 - effortless
- Cons:
 - decreasing resource usage

Storing Data into Local Filesystem of the Specified Mesos Slave

- Scenarios:
 - Data disk is NOT included into Mesos resource
 - Avoid of data revoked by Mesos
 - The next task can restore data
- Pros:
 - Using the CPU, RAM, network resource of the Mesos cluster: one single dockerized MySQL task
- Cons:
 - Resource STATIC RESERVED to make sure the stateful service can be launched on the specified slave successfully always
 - Mesos can NOT release the outdated persisted data, have to manually do it
 - Data conflict: avoid it by launching one task per slave temporarily
 - how to share the data for HA?

Setting DFS to get rid of static reserve

- Pros:
 - HA
 - Get rid of static reserve
- Cons:
 - network delay of data transfer. that can be accepted by MySQL/PG?
 - The network resource taken by data transfer is OUT of CONTROL: Setting another data transfer network to avoid it? => decreasing the network resource usage again

Taking advantage of stateful service build-in data-replication feature

- Scenarios:
 - Multiple stateful services are supporting data-replication: Cassandra, MariaDB Galera, MongoDB, etc.
 - Cassandra on Mesos: <https://github.com/mesosphere/cassandra-mesos>
 - MariaDB Galera on Mesos: <http://sttts.github.io/galera/mesos/2015/03/04/galera-on-mesos.html>
- Pros:
 - Network is in control via task build-in data replication
 - Stateful service itself is responsible for network delay
- Cons:
 - Persistent disk is outside of cluster still

in a nutshell

To persistent data, we
have to store data outside
of the cluster currently

Disk Isolation and Monitoring

- To keep other tasks running, Mesos is restricting disk quota per task

Mesos disk is a GENERIC disk

- Separated Filesystem
 - Created by physical disk/LVM etc.
 - Signal ENOSPC triggered once data size expanded
 - Task is terminated
 - Hard enforcement
 - Adapt to production environment
- Shared Filesystem
 - Sharing file directory with other tasks
 - Monitoring “disk” usage by executing “du” periodically
 - Tolerating data size expansion
 - Soft enforcement

Shared Filesystem needs disk isolator

- Support Mesos build-in container only.
- Map container path to slave host path by command “mount -n --bind”. Kernel will umount it automatically after task finished, at the same time.
- Use docker volume mapper for docker container

Disk IO resource

- Heavy-disk-IO task is throttling other tasks maybe
- Try “Cgroups blkio controller” for disk IO isolation in future

Persistent Volumes

- Data in persistent volumes WON'T be GC after task completed
- Another task can restore last finished task data
- Data survives even slave info/id changed or rebooted
- Belonging to Mesos cluster resource, handled by Mesos cluster

Regular Resource VS. Persistent Resource

- Regular resource
 - is renamed from current 0.22 resource
 - match stateless task
 - CPU, RAM, Disk will be GC after task completed
- Persistent resource
 - Besides persistent volumes, CPU, RAM is included also. WHY?
 - match stateful service
 - data in persistent volume is reserved after task completed
 - be validated by google Borg

Resource Offer with persistent resource

```
{ "id" : { "value" : "offerid-123456789" },  
  "framework_id" : "MYID",  
  "slave_id" : "slaveid-123456789",  
  "hostname" : "hostname1.prod.twtr.net"  
  "resources" : [  
    // Regular resources.  
    { "name" : "cpu", "type" : SCALAR, "scalar" : 10 }  
    { "name" : "mem", "type" : SCALAR, "scalar" : 12GB }  
    { "name" : "disk", "type" : SCALAR, "scalar" : 90GB }  
    // Persistent resources.  
    { "name" : "cpu", "type" : SCALAR, "scalar" : 2,  
      "persistence" : { "framework_id" : "MYID", "handle" : "uuid-123456789" } }  
    { "name" : "mem", "type" : SCALAR, "scalar" : 2GB,  
      "persistence" : { "framework_id" : "MYID", "handle" : "uuid-123456789" } }  
    { "name" : "disk", "type" : SCALAR, "scalar" : 10GB,  
      "persistence" : { "framework_id" : "MYID", "handle" : "uuid-123456789" } }  
  ]  
  ...  
}
```


Shared persistent data

- MySQL framework maybe launch 2 tasks accessing the same persistent data, for example
 - mysqld server
 - data backup periodically
- under discussion still

Dynamic Reservations

- Why reservations? Make sure specified framework running on specified slave(s)
- What reserved? CPU/RAM/DISK/Network
- How to reserve? Static reserve by setting slave role currently

Dynamic Reservations

- Framework can broadcast reserving resource dynamically, whatever regular or persistent resource, by setting “reserved_role”, when launch a task
- The “reserved_role” can be reoffered to the same framework after task completed
- Framework can release the dynamic reserved resource by itself, while can NOT release the before static one
- Very different from the before static reserve by setting slave role

What we can contribute for Mesos persistent Storage

- acentric or node-equal stateful service
 - Cassandra, for example
 - effortless: dockerize cassandra program => distribute to the multiple specified slaves over scheduler
- master-slave or leader-follower stateful service
 - HDFS, MongoDB, MySQL, for example
 - name node, data node, config node
 - develop different framework for different service
 - Let the framework solve fault-tolerant, backup issue

What we can contribute for Mesos persistent Storage

- know more the stateful service content
 - mail-list
 - source code
 - doc
- try to develop the framework or dockerize the program
- solve the fault-tolerate, backup issue
- avoid to re-invent wheel

Q&A

We are hiring

- Python
- Django
- AngularJS
- Mesos
- Docker
- Linux
- Git
- opensource
- shell
- 3+ years
- 望京
- line 14/15 subway
- work with googler, red hatter, 活力四射er
- more geek, more money

<mailto:jjyan@dataman-inc.com>

Thanks